

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from scipy.stats import pearsonr
import seaborn as sns

data = pd.read_csv(r"C:/Users/Nishi/Downloads/diabetic_data-1.csv")
```

```
In [2]: data.shape
```

```
Out[2]: (101766, 50)
```

```
In [3]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 101766 entries, 0 to 101765
Data columns (total 50 columns):
```

#	Column	Non-Null Count	Dtype
0	encounter_id	101766 non-null	int64
1	patient_nbr	101766 non-null	int64
2	race	101766 non-null	object
3	gender	101766 non-null	object
4	age	101766 non-null	object
5	weight	101766 non-null	object
6	admission_type_id	101766 non-null	int64
7	discharge_disposition_id	101766 non-null	int64
8	admission_source_id	101766 non-null	int64
9	time_in_hospital	101766 non-null	int64
10	payer_code	101766 non-null	object
11	medical_specialty	101766 non-null	object
12	num_lab_procedures	101766 non-null	int64
13	num_procedures	101766 non-null	int64
14	num_medications	101766 non-null	int64
15	number_outpatient	101766 non-null	int64
16	number_emergency	101766 non-null	int64
17	number_inpatient	101766 non-null	int64
18	diag_1	101766 non-null	object
19	diag_2	101766 non-null	object
20	diag_3	101766 non-null	object
21	number_diagnoses	101766 non-null	int64
22	max_glu_serum	5346 non-null	object
23	A1Cresult	17018 non-null	object
24	metformin	101766 non-null	object
25	repaglinide	101766 non-null	object
26	nateglinide	101766 non-null	object
27	chlorpropamide	101766 non-null	object
28	glimepiride	101766 non-null	object
29	acetohexamide	101766 non-null	object
30	glipizide	101766 non-null	object
31	glyburide	101766 non-null	object
32	tolbutamide	101766 non-null	object
33	pioglitazone	101766 non-null	object
34	rosiglitazone	101766 non-null	object
35	acarbose	101766 non-null	object
36	miglitol	101766 non-null	object
37	troglitazone	101766 non-null	object
38	tolazamide	101766 non-null	object
39	examide	101766 non-null	object
40	citoglipton	101766 non-null	object
41	insulin	101766 non-null	object
42	glyburide-metformin	101766 non-null	object
43	glipizide-metformin	101766 non-null	object
44	glimepiride-pioglitazone	101766 non-null	object
45	metformin-rosiglitazone	101766 non-null	object
46	metformin-pioglitazone	101766 non-null	object
47	change	101766 non-null	object
48	diabetesMed	101766 non-null	object
49	readmitted	101766 non-null	object

```
dtypes: int64(13), object(37)
```

```
memory usage: 38.8+ MB
```

```
In [4]: nan_counts = data.isna().sum()  
print(nan_counts)
```

encounter_id	0
patient_nbr	0
race	0
gender	0
age	0
weight	0
admission_type_id	0
discharge_disposition_id	0
admission_source_id	0
time_in_hospital	0
payer_code	0
medical_specialty	0
num_lab_procedures	0
num_procedures	0
num_medications	0
number_outpatient	0
number_emergency	0
number_inpatient	0
diag_1	0
diag_2	0
diag_3	0
number_diagnoses	0
max_glu_serum	96420
A1Cresult	84748
metformin	0
repaglinide	0
nateglinide	0
chlorpropamide	0
glimepiride	0
acetohexamide	0
glipizide	0
glyburide	0
tolbutamide	0
pioglitazone	0
rosiglitazone	0
acarbose	0
miglitol	0
troglitazone	0
tolazamide	0
examide	0
citoglipton	0
insulin	0
glyburide-metformin	0
glipizide-metformin	0
glimepiride-pioglitazone	0
metformin-rosiglitazone	0
metformin-pioglitazone	0
change	0
diabetesMed	0
readmitted	0
dtype: int64	

```
In [5]: data.max_glu_serum.value_counts()
```

```
Out[5]: max_glu_serum
      Norm      2597
      >200     1485
      >300     1264
      Name: count, dtype: int64
```

```
In [6]: data.max_glu_serum.mode()
```

```
Out[6]: 0      Norm
      Name: max_glu_serum, dtype: object
```

```
In [7]: data.A1Cresult.value_counts()
```

```
Out[7]: A1Cresult
      >8      8216
      Norm   4990
      >7      3812
      Name: count, dtype: int64
```

```
In [8]: data.head(20)
```

Out[8]:

	encounter_id	patient_nbr	race	gender	age	weight	admission_type_id	d
0	2278392	8222157	Caucasian	Female	[0-10)	?		6
1	149190	55629189	Caucasian	Female	[10-20)	?		1
2	64410	86047875	AfricanAmerican	Female	[20-30)	?		1
3	500364	82442376	Caucasian	Male	[30-40)	?		1
4	16680	42519267	Caucasian	Male	[40-50)	?		1
5	35754	82637451	Caucasian	Male	[50-60)	?		2
6	55842	84259809	Caucasian	Male	[60-70)	?		3
7	63768	114882984	Caucasian	Male	[70-80)	?		1
8	12522	48330783	Caucasian	Female	[80-90)	?		2
9	15738	63555939	Caucasian	Female	[90-100)	?		3
10	28236	89869032	AfricanAmerican	Female	[40-50)	?		1
11	36900	77391171	AfricanAmerican	Male	[60-70)	?		2
12	40926	85504905	Caucasian	Female	[40-50)	?		1
13	42570	77586282	Caucasian	Male	[80-90)	?		1
14	62256	49726791	AfricanAmerican	Female	[60-70)	?		3
15	73578	86328819	AfricanAmerican	Male	[60-70)	?		1
16	77076	92519352	AfricanAmerican	Male	[50-60)	?		1
17	84222	108662661	Caucasian	Female	[50-60)	?		1
18	89682	107389323	AfricanAmerican	Male	[70-80)	?		1

	encounter_id	patient_nbr	race	gender	age	weight	admission_type_id	d
19	148530	69422211	?	Male	[70-80)	?	3	

20 rows × 50 columns

```
In [9]: data.weight.value_counts()
```

```
Out[9]: weight
?          98569
[75-100)    1336
[50-75)      897
[100-125)    625
[125-150)    145
[25-50)       97
[0-25)       48
[150-175)    35
[175-200)    11
>200         3
Name: count, dtype: int64
```

```
In [10]: data.drop('max_glu_serum', axis=1, inplace=True)
data.drop('A1Cresult', axis=1, inplace=True)
data.drop('weight', axis=1, inplace=True)
```

```
In [11]: data.age.value_counts()
```

```
Out[11]: age
[70-80)    26068
[60-70)    22483
[50-60)    17256
[80-90)    17197
[40-50)     9685
[30-40)     3775
[90-100)    2793
[20-30)    1657
[10-20)     691
[0-10)     161
Name: count, dtype: int64
```

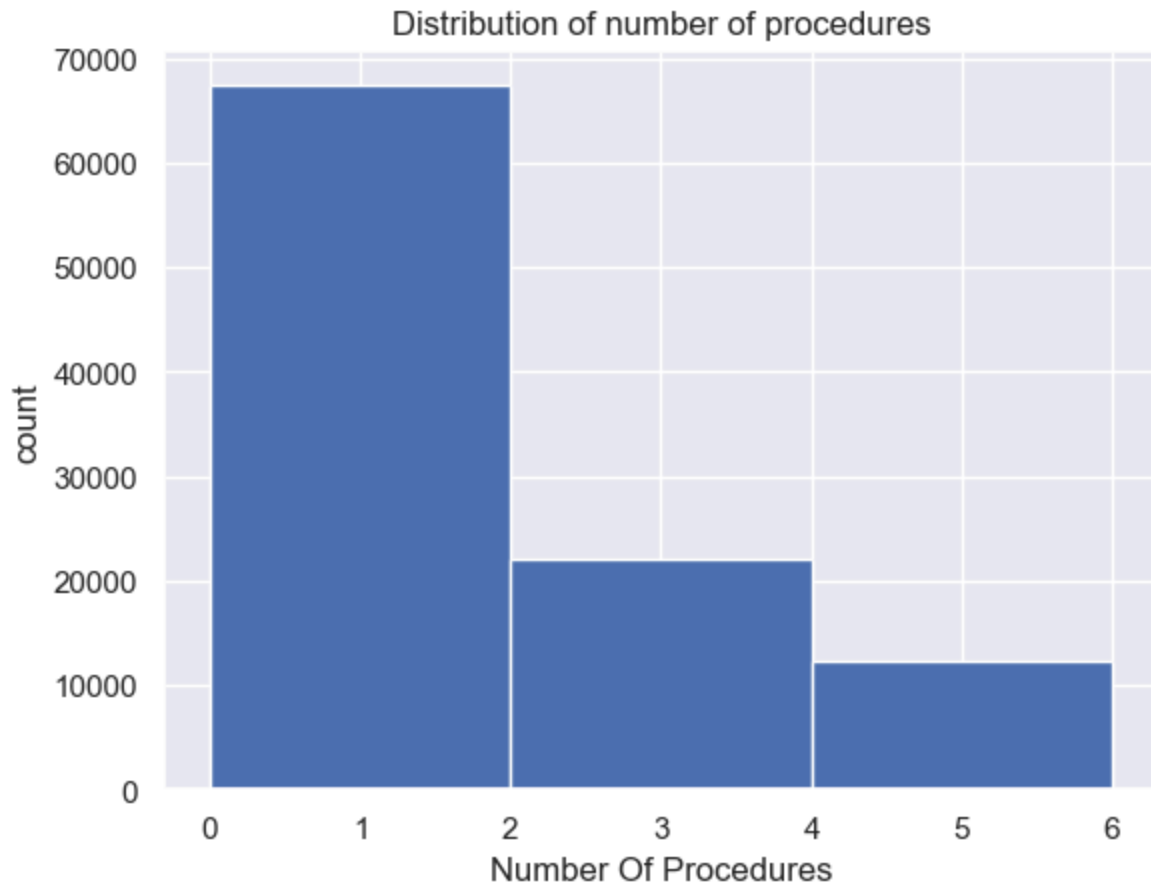
```
In [12]: data.time_in_hospital.value_counts()
```

```
Out[12]: time_in_hospital
3      17756
2      17224
1      14208
4      13924
5       9966
6       7539
7       5859
8       4391
9       3002
10      2342
11      1855
12      1448
13      1210
14       1042
Name: count, dtype: int64
```

```
In [13]: data.num_procedures.value_counts()
```

```
Out[13]: num_procedures
0      46652
1      20742
2      12717
3       9443
6       4954
4       4180
5       3078
Name: count, dtype: int64
```

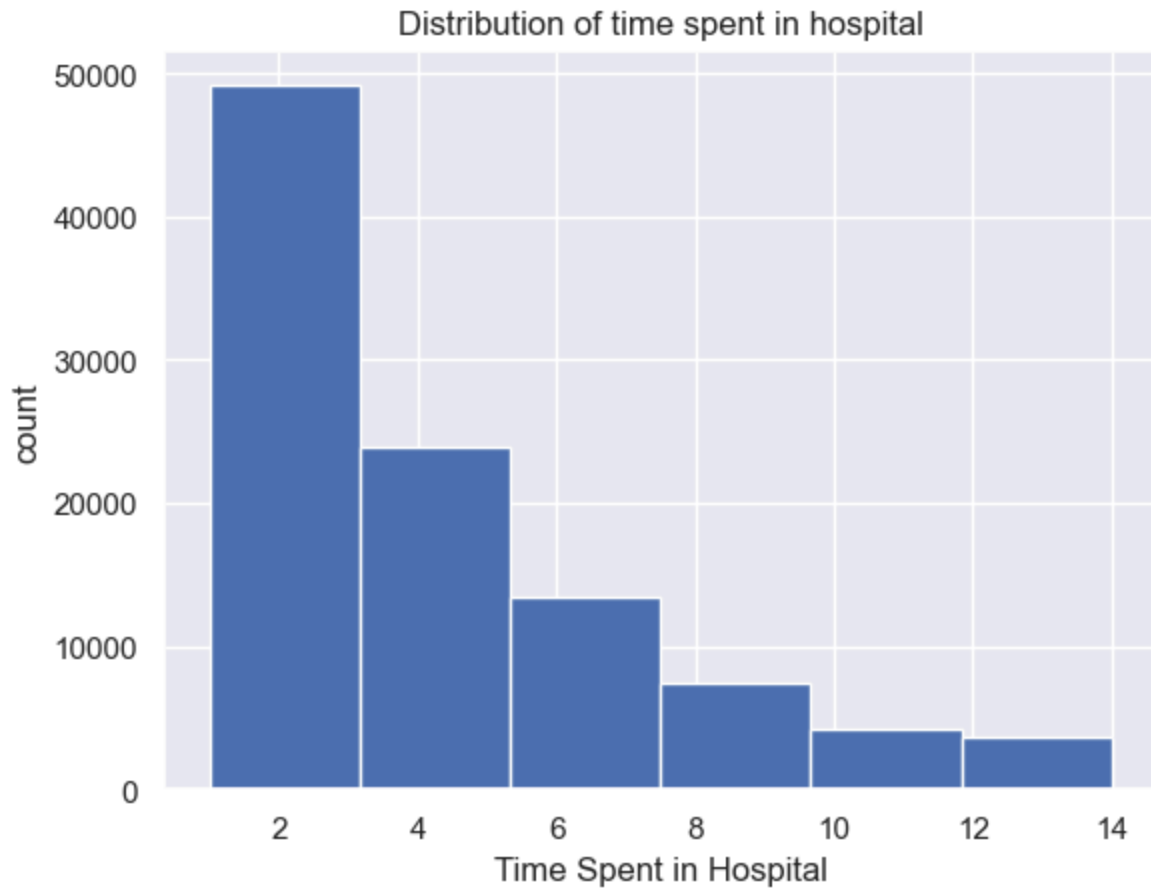
```
In [14]: sns.set()
num_of_procs = data.num_procedures
plt.hist(num_of_procs, bins=3)
# Label axes
plt.xlabel('Number Of Procedures')
plt.ylabel('count')
plt.title('Distribution of number of procedures')
# Show histogram
plt.show()
```


```
In [15]: data.time_in_hospital.value_counts()
```

```
Out[15]: time_in_hospital
3      17756
2      17224
1      14208
4      13924
5       9966
6       7539
7       5859
8       4391
9       3002
10      2342
11      1855
12      1448
13      1210
14       1042
Name: count, dtype: int64
```

```
In [16]: sns.set()
time_in_hosp = data.time_in_hospital
plt.hist(time_in_hosp, bins=6)
# Label axes
plt.xlabel('Time Spent in Hospital')
plt.ylabel('count')
plt.title('Distribution of time spent in hospital')
# Show histogram
plt.show()
```



```
In [17]: sns.set()
Age = data.age
plt.hist(Age, bins=5)
# Label axes
plt.xlabel('Age')
plt.ylabel('count')
plt.title('Distribution of age')
plt.xticks(rotation=45)
# Show histogram
plt.show()
```



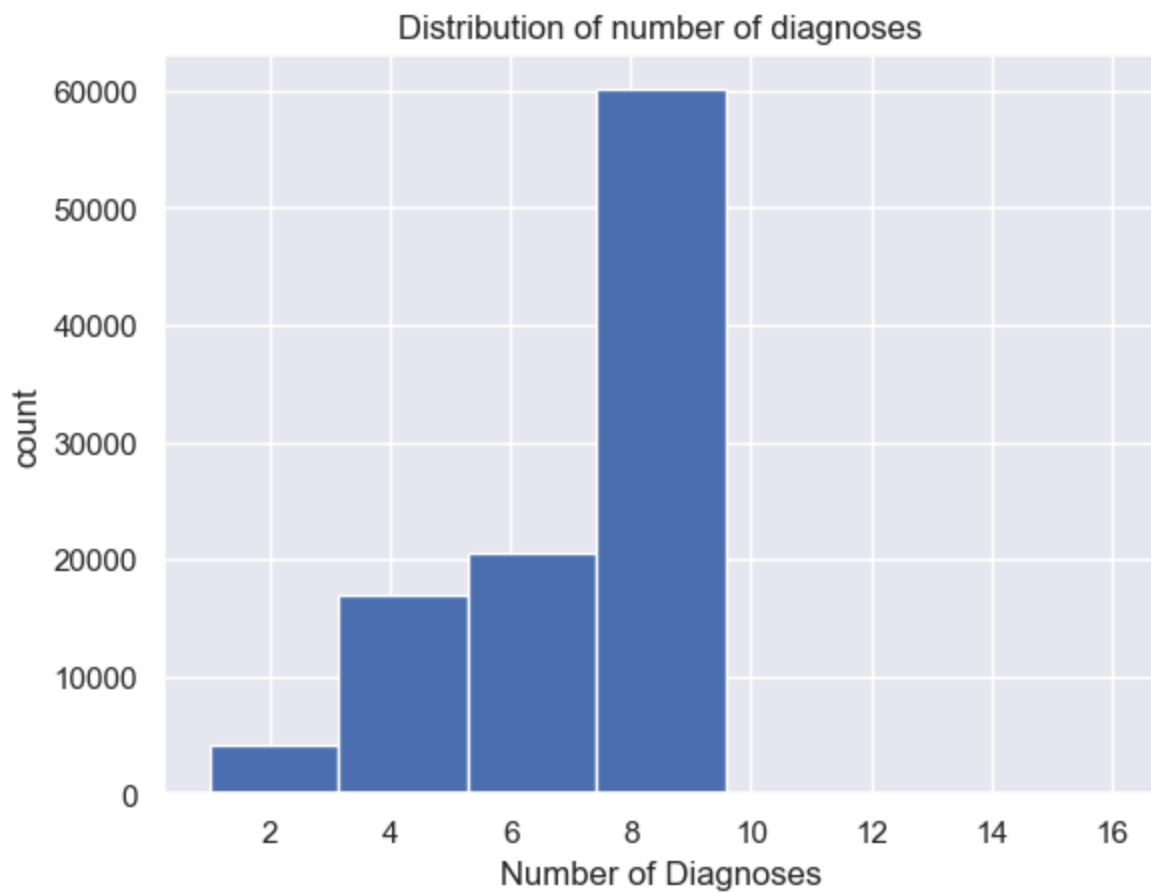
```
In [18]: data.age.value_counts()
```

```
Out[18]: age
[70-80)    26068
[60-70)    22483
[50-60)    17256
[80-90)    17197
[40-50)     9685
[30-40)     3775
[90-100)     2793
[20-30)     1657
[10-20)      691
[0-10)       161
Name: count, dtype: int64
```

```
In [19]: data.number_diagnoses.value_counts()
```

```
Out[19]: number_diagnoses
9      49474
5      11393
8      10616
7      10393
6      10161
4       5537
3       2835
2       1023
1        219
16        45
10         17
13         16
11          11
15          10
12           9
14           7
Name: count, dtype: int64
```

```
In [20]: sns.set()
diagnoses = data.number_diagnoses
plt.hist(diagnoses, bins=7)
# Label axes
plt.xlabel('Number of Diagnoses')
plt.ylabel('count')
plt.title('Distribution of number of diagnoses')
# Show histogram
plt.show()
```

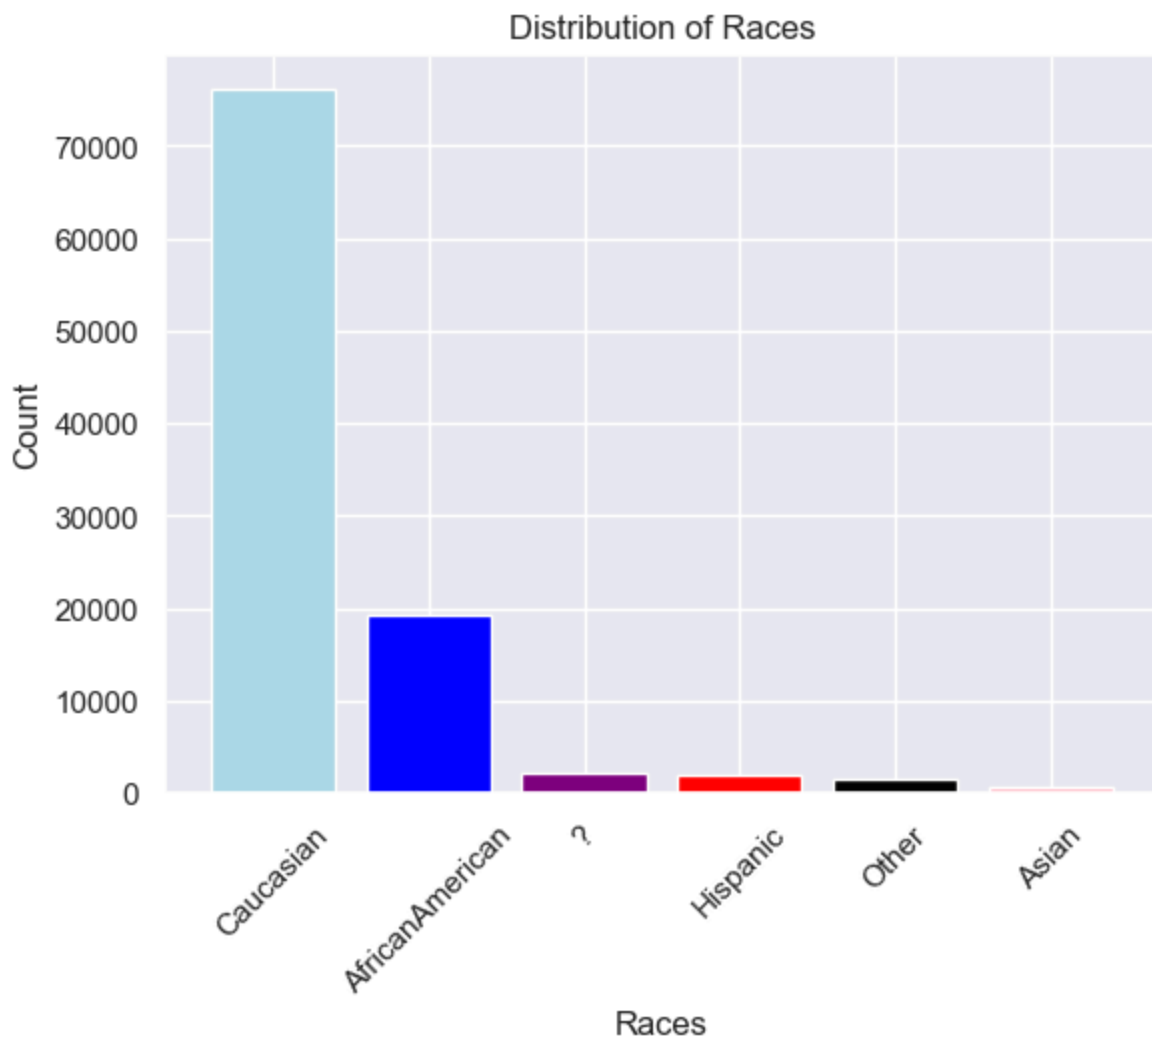


```
In [21]: data.readmitted.value_counts()
```


```
Out[21]: readmitted  
NO      54864  
>30    35545  
<30     11357  
Name: count, dtype: int64
```

```
In [22]: race_counts = data.race.value_counts()  
gender_counts = data.gender.value_counts()
```

```
In [23]: # creating vertical bar chart  
fig, ax = plt.subplots()  
color = ['lightblue', 'blue', 'purple', 'red', 'black', 'pink']  
ax.bar(race_counts.keys(), race_counts.values, color=color)  
  
plt.xticks(rotation=45)  
ax.set_ylabel('Count')  
ax.set_xlabel('Races')  
ax.set_title('Distribution of Races')  
plt.show()
```



```
In [24]: fig, ax = plt.subplots()
ax.pie(gender_counts.values, labels=gender_counts.keys(), autopct='%1.1f%%', startangle=90)
ax.axis('equal')
plt.title('Gender Distribution')
plt.show()
```



```
In [25]: data.diabetesMed.value_counts()
```

```
Out[25]: diabetesMed
Yes      78363
No       23403
Name: count, dtype: int64
```

```
In [26]: data.insulin.value_counts()
```


```
Out[26]: insulin
No      47383
Steady  30849
Down    12218
Up      11316
Name: count, dtype: int64
```

```
In [27]: diabetesMedCount = data.diabetesMed.value_counts()
insulinCount = data.insulin.value_counts()
```

```
In [28]: # creating vertical bar chart
fig, ax = plt.subplots()
color = ['lightblue', 'blue', 'purple', 'red', 'black', 'pink']
ax.bar(insulinCount.keys(), insulinCount.values, color=color)
```


```
plt.xticks(rotation=45)
ax.set_ylabel('Count')
ax.set_xlabel('insulin')

plt.show()
```



```
In [29]: fig, ax = plt.subplots()
ax.pie(diabetesMedCount.values, labels=diabetesMedCount.keys(), autopct='%1.1f%%',
ax.axis('equal')
plt.title('Diabetes Med Count Distribution')
plt.show()
```

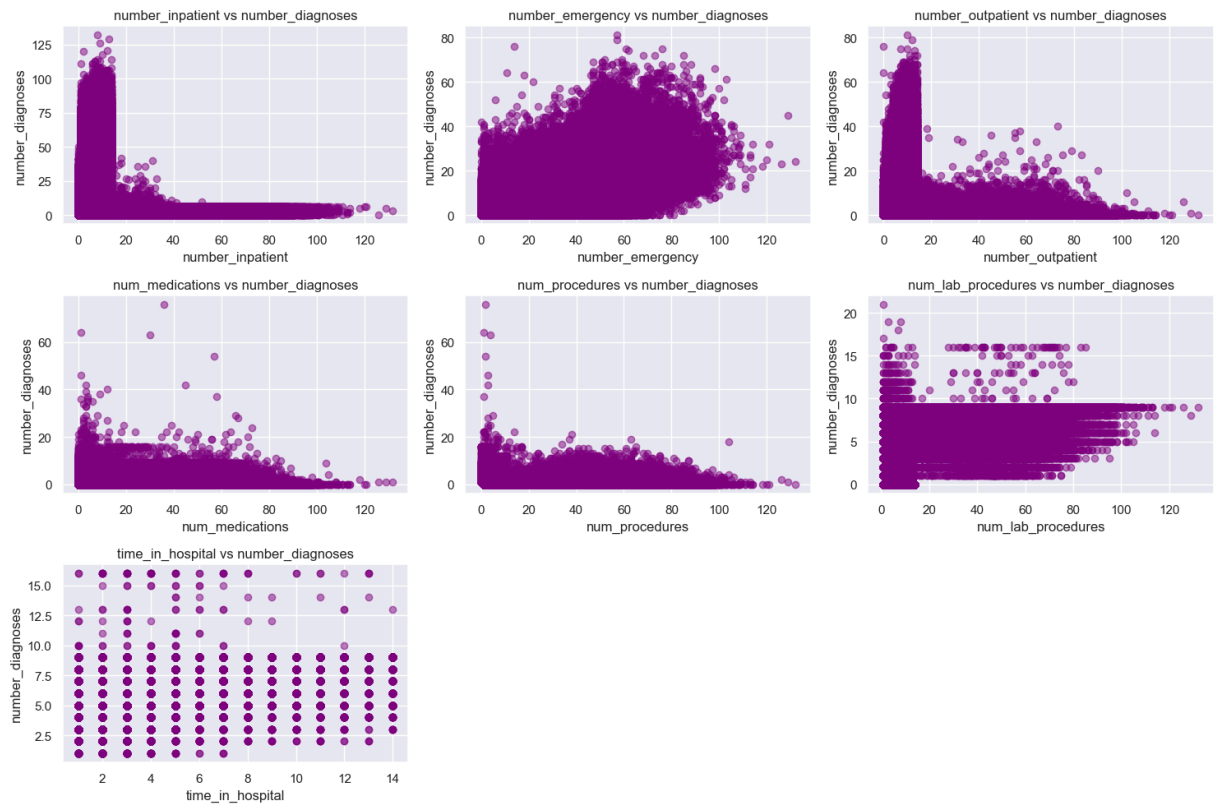
Diabetes Med Count Distribution



```
In [30]: import matplotlib.pyplot as plt

continuous_vars = ['time_in_hospital', 'num_lab_procedures', 'num_procedures',
                  'num_medications', 'number_outpatient', 'number_emergency',
                  'number_inpatient', 'number_diagnoses']


# Generate scatter plots for each pair of continuous variables
plt.figure(figsize=(15, 10))
for i in range(len(continuous_vars)):
    for j in range(i + 1, len(continuous_vars)):
        plt.subplot(3, 3, j - i)
        plt.scatter(data[continuous_vars[i]], data[continuous_vars[j]], color='purple')
        plt.xlabel(continuous_vars[i])
        plt.ylabel(continuous_vars[j])
        plt.title(f'{continuous_vars[i]} vs {continuous_vars[j]}')
plt.tight_layout()
plt.show()
```

```
In [31]: numerical_data = data.select_dtypes(include=[np.number])

# Calculating the correlation matrix
full_correlation_matrix = numerical_data.corr()

# Creating a heatmap to visualize the full correlation matrix
plt.figure(figsize=(14, 12))
full_heatmap = sns.heatmap(full_correlation_matrix, annot=True, cmap='coolwarm', fm
plt.title('Heatmap of Correlation Matrix')
plt.show()
```




```
In [32]: # Create a catplot to visualize the relationship between race, insulin usage, and
sns.set(style="whitegrid")

# Create the plot
g = sns.catplot(
    data=data,
    x="race",
    hue="diabetesMed",
    col="insulin",
    kind="count",
    height=5,
    aspect=1.5,
    palette="viridis"
)


# Set the titles and Labels
g.set_titles(col_template="{col_name} insulin usage")
g.set_axis_labels("Race", "Count")
g._legend.set_title('Diabetes Medication')

plt.xticks(rotation=45)
```

```
plt.tight_layout()
plt.show()
```



```
In [33]: plt.figure(figsize=(12, 8))
box_plot = sns.boxplot(x='variable', y='value', data=pd.melt(data[['time_in_hospital',
box_plot.set_title('Box Plot of Numerical Variables')
box_plot.set_xlabel('Variable')
box_plot.set_ylabel('Value')
plt.xticks(rotation=45) # Rotate labels for better readability
plt.show()
```



```
In [35]: # Columns to be dropped
columns_to_drop = [
    'encounter_id',
    'patient_nbr',
    'admission_type_id',
    'discharge_disposition_id',
    'admission_source_id',
    'payer_code'
]
```

```
data2 = data.drop(columns=columns_to_drop)
```

```
In [36]: #Remove outliers

# Define a function to remove outliers using the IQR method
def remove_outliers(df, columns):
    for column in columns:
        Q1 = df[column].quantile(0.25)
        Q3 = df[column].quantile(0.75)
        IQR = Q3 - Q1
        lower_bound = Q1 - 1.5 * IQR
        upper_bound = Q3 + 1.5 * IQR
        df = df[(df[column] >= lower_bound) & (df[column] <= upper_bound)]
    return df

# List of numerical columns to check for outliers
numerical_columns = [
    'time_in_hospital', 'num_lab_procedures', 'num_procedures',
    'num_medications', 'number_outpatient', 'number_emergency',
    'number_inpatient', 'number_diagnoses'
]

# Remove outliers from the dataset
data2 = remove_outliers(data2, numerical_columns)

print(data2.head())
```

	race	gender	age	time_in_hospital	medical_specialty	\
1	Caucasian	Female	[10-20)	3	?	
3	Caucasian	Male	[30-40)	2	?	
4	Caucasian	Male	[40-50)	1	?	
6	Caucasian	Male	[60-70)	4	?	
7	Caucasian	Male	[70-80)	5	?	

	num_lab_procedures	num_procedures	num_medications	number_outpatient	\
1	59	0	18	0	
3	44	1	16	0	
4	51	0	8	0	
6	70	1	21	0	
7	73	0	12	0	

	number_emergency	...	citoglipton	insulin	glyburide-metformin	\
1	0	...	No	Up	No	
3	0	...	No	Up	No	
4	0	...	No	Steady	No	
6	0	...	No	Steady	No	
7	0	...	No	No	No	

	glipizide-metformin	glimepiride-pioglitazone	metformin-rosiglitazone	\
1	No		No	No
3	No		No	No
4	No		No	No
6	No		No	No
7	No		No	No

	metformin-pioglitazone	change	diabetesMed	readmitted
1	No	Ch	Yes	>30
3	No	Ch	Yes	NO
4	No	Ch	Yes	NO
6	No	Ch	Yes	NO
7	No	No	Yes	>30

[5 rows x 41 columns]

In [37]: data2.readmitted.value_counts()

Out[37]: readmitted

NO	39280
>30	20943
<30	6369

Name: count, dtype: int64

In [38]: data2['readmitted'] = data2['readmitted'].apply(lambda x: 'YES' if x == '<30' else data2.readmitted.value_counts())

Out[38]: readmitted

NO	60223
YES	6369

Name: count, dtype: int64

In [40]: from sklearn.preprocessing import MinMaxScaler, LabelEncoder
 from sklearn.model_selection import train_test_split
 from sklearn.tree import DecisionTreeClassifier

```

from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import confusion_matrix
from imblearn.over_sampling import SMOTE

# Normalize continuous variables
scaler = MinMaxScaler()
continuous_vars = ['time_in_hospital', 'num_lab_procedures', 'num_procedures',
                  'num_medications', 'number_outpatient', 'number_emergency',
                  'number_inpatient', 'number_diagnoses']
data2[continuous_vars] = scaler.fit_transform(data2[continuous_vars])

# Check for and handle missing values
data2.fillna(method='ffill', inplace=True)

# Encode categorical variables and target variable
categorical_cols = data2.select_dtypes(include=['object']).columns.tolist()
# Remove the target column 'readmitted' from categorical columns
categorical_cols.remove('readmitted')

# One-hot encoding categorical variables
data_encoded2 = pd.get_dummies(data2, columns=categorical_cols, drop_first=True)

# Encoding the target variable
label_encoder = LabelEncoder()
data_encoded2['readmitted_encoded'] = label_encoder.fit_transform(data2['readmitted'])

# Splitting the data
X = data_encoded2.drop(['readmitted', 'readmitted_encoded'], axis=1) # Features
y = data_encoded2['readmitted_encoded'] # Target variable

# Apply SMOTE to balance the dataset
smote = SMOTE(random_state=42)
X_balanced, y_balanced = smote.fit_resample(X, y)

X_train, X_test, y_train, y_test = train_test_split(X_balanced, y_balanced, test_si

```

C:\Users\Nishi\AppData\Local\Temp\ipykernel_12368\1867024009.py:16: FutureWarning: DataFrame.fillna with 'method' is deprecated and will raise in a future version. Use obj.ffill() or obj.bfill() instead.

```
data2.fillna(method='ffill', inplace=True)
```

```

In [44]: from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import confusion_matrix, precision_score, recall_score, f1_score

# Fitting the DecisionTreeClassifier
DecisionTree = DecisionTreeClassifier(random_state=0, max_depth=9, min_samples_leaf=1)
DecisionTree.fit(X_train, y_train)

# Evaluate the DecisionTreeClassifier model
y_pred = DecisionTree.predict(X_test)
conf_matrix = confusion_matrix(y_test, y_pred)
print("Confusion Matrix (Decision Tree):")
print(conf_matrix)

# Calculate metrics

```

```

accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred, pos_label=1)
recall = recall_score(y_test, y_pred, pos_label=1)
f1 = f1_score(y_test, y_pred, pos_label=1)

# Print out the metrics
print(f"Accuracy: {accuracy:.4f}")
print(f"Precision (yes): {precision:.4f}")
print(f"Recall (yes): {recall:.4f}")
print(f"F1-score (yes): {f1:.4f}")

```

Confusion Matrix (Decision Tree):

```

[[11757   190]
 [ 3569  8574]]

```

Accuracy: 0.8440

Precision (yes): 0.9783

Recall (yes): 0.7061

F1-score (yes): 0.8202

In [45]: `from sklearn.ensemble import RandomForestClassifier`
`from sklearn.metrics import confusion_matrix, precision_score, recall_score, f1_score`

```

# Fitting the RandomForestClassifier
rf_classifier = RandomForestClassifier(random_state=0)
rf_classifier.fit(X_train, y_train)

# Evaluate the RandomForestClassifier model
y_pred = rf_classifier.predict(X_test)
conf_matrix = confusion_matrix(y_test, y_pred)
print("Confusion Matrix (Random Forest):")
print(conf_matrix)

# Calculate metrics
accuracy = accuracy_score(y_test, y_pred)
precision_yes = precision_score(y_test, y_pred, pos_label=1)
recall_yes = recall_score(y_test, y_pred, pos_label=1)
f1_yes = f1_score(y_test, y_pred, pos_label=1)

# Print out the metrics
print(f"Accuracy: {accuracy:.4f}")
print(f"Precision (yes): {precision_yes:.4f}")
print(f"Recall (yes): {recall_yes:.4f}")
print(f"F1-score (yes): {f1_yes:.4f}")

```

Confusion Matrix (Random Forest):

```

[[11910    37]
 [  684 11459]]

```

Accuracy: 0.9701

Precision (yes): 0.9968

Recall (yes): 0.9437

F1-score (yes): 0.9695

In [46]: `from sklearn.ensemble import GradientBoostingClassifier`
`from sklearn.metrics import confusion_matrix, precision_score, recall_score, f1_score`

```

# Fitting the GradientBoostingClassifier
gb_classifier = GradientBoostingClassifier(random_state=0)

```

```
gb_classifier.fit(X_train, y_train)

# Evaluate the GradientBoostingClassifier model
y_pred = gb_classifier.predict(X_test)
conf_matrix = confusion_matrix(y_test, y_pred)
print("Confusion Matrix (Gradient Boosting):")
print(conf_matrix)

# Calculate metrics
accuracy = accuracy_score(y_test, y_pred)
precision_yes = precision_score(y_test, y_pred, pos_label=1)
recall_yes = recall_score(y_test, y_pred, pos_label=1)
f1_yes = f1_score(y_test, y_pred, pos_label=1)

# Print out the metrics
print(f"Accuracy: {accuracy:.4f}")
print(f"Precision (yes): {precision_yes:.4f}")
print(f"Recall (yes): {recall_yes:.4f}")
print(f"F1-score (yes): {f1_yes:.4f}")
```

Confusion Matrix (Gradient Boosting):

```
[[11945    2]
 [ 1966 10177]]
```

Accuracy: 0.9183

Precision (yes): 0.9998

Recall (yes): 0.8381

F1-score (yes): 0.9118

In []: